

G54MDP


Mobile Device Programming

Lecture 18 – Power and Batteries




Android Power Management

- Built as a wrapper around Linux Power Management
- In the kernel
 - Added **Early Suspend** mechanism
 - Added **Partial Wake Lock** mechanism
- Apps and services must request CPU resource in order to keep power on
 - Otherwise Android will shut down the CPU
 - Suspend operational RAM to NAND
- Wake locks and timeouts constantly switch the state of the system's power
 - Overall system power consumption decreases
 - “Better” use of battery capacity



Suspended Android

- Running applications / services are suspended
- CPU is powered down
 - Phone is not off
- Other components (SOC) continue to operate
 - CPU is periodically woken to handle scheduled tasks
 - Real time clock manifests as /dev/alarm
 - AlarmManager – Alarms, email polling...
 - GSM modem will wake CPU on call / SMS notifications
- Why use a PARTIAL_WAKE_LOCK?
 - Playing music – does not require screen to be on
 - Avoid suspension during periodic tasks
 - Android will try to suspend even when it is checking whether the alarm clock should sound
 - AlarmManager acquires, then releases a PARTIAL_WAKE_LOCK



Application Wake Locks

- Provides user-space (application) ability to manage power indirectly
 - Request a wake lock
- Application flow
 - Acquire a handle to the static `PowerManager` service with `Context.getSystemService()`
 - Create a wake lock and specify flags for screen, backlight etc
 - Acquire the wake lock
 - Perform the operation
 - Play MP3
 - Release the wake lock
- Must be used carefully
 - Keeping a wake lock for a long period of time will trash battery life
 - The CPU will not be allowed to sleep
- Tasks scheduled using the `AlarmManager` do not require a wake lock
 - `AlarmManager` acquires the lock while calling our scheduled task

Kernel Wake Locks

- Used to prevent the system entering suspended mode
 - Can be acquired and released by native code, or directly from within the kernel
 - Partial Wake Locks all reside in the kernel as they keep the CPU processing
- A single kernel wake lock manages multiple user mode (java) wake locks
 - PowerManagerService native kernel code partial wake lock
 - Audio driver partial wake lock while playing audio
 - Kernel has one last partial wake lock that exists to keep the kernel alive while other wake locks exist

WakeLock in
Java layer



FULL_WAKE_LOCK

SCREEN_BRIGHT_WAKE_LOCK

PARTIAL_WAKE_LOCK

SCREEN_DIM_WAKE_LOCK

WakeLock
in Kernel



"PowerManagerService"
PARTIAL_WAKE_LOCK

"main"
PARTIAL_WAKE_LOCK

PARTIAL_WAKE_LOCK

Acquiring a Wake Lock

- Request sent to PowerManager (java) to acquire a wake lock
- PowerManagerService notified to take a wake lock
 - Add wake lock to an internal list
 - Set the requested power state
 - If this is the first partial wake lock take a kernel partial wake lock
 - This will protect all the partial wake locks
 - For subsequent wake locks simply add to the list

Releasing a Wake Lock

- Request sent to PowerManager (java) to release the wake lock
- Wake lock removed from the internal list
- If the wake lock is the last partial wake lock in the list
 - Release the kernel wake lock
- If kernel main wake lock is the only wake lock
 - Release main kernel wake lock
 - Device moves to suspend

Early Suspend / Late Resume

- More modifications to the Linux kernel
- In standard Linux all modules are suspended / resumed at the same time
 - Suspend
 - Freeze all user processes and kernel tasks
 - Call the suspend function for all devices
 - Suspend the kernel and suspend the CPU
 - Resume
 - Wake up the kernel
 - Wake up the registered devices
 - Unfreeze user processes and resume kernel tasks

Early Suspend / Late Resume

- Suspend as much as possible even if the kernel is still operating
- Early suspend
 - **Between** screen-off and full suspension
 - Tells devices to attempt to suspend even though a wake lock may be keeping the kernel awake
 - Stop screen, touch screen, backlight, close drivers
 - **Note** difference between “screen is on” and “kernel screen device is awake”!
- Cannot achieve full suspension (stop CPU, RAM -> NAND) until all wake locks are released
 - However attempts to suspend as much as possible
- Late resume
 - Kernel devices that were early_suspended are subsequently late_resumed
 - Can wake the kernel without waking up the entire device
 - Resume suspended devices once the kernel is awake and working

AlarmManager

- Schedule an application to be run at some point in the future
 - As usual, specify an Intent to be broadcast at some time
 - At a regular interval, after an elapsed time
 - Extend Broadcast Receiver
 - onReceive method is called when the alarm goes off
- AlarmManager is triggered regularly by the real time clock device in the kernel
 - Alarms go off even when the CPU is asleep
 - AlarmManager holds a partial wake lock while onReceive is executing
 - If the alarm starts a service, need to acquire our own partial wake lock, as the system may go to sleep while the service is starting
 - Default off

Coding for Battery Life

- Android constantly tries to suspend
 - Extends battery life on average
- Apps allow us to **do things** with the phone
 - Use the CPU, use the radio / network
 - reduce battery life
- Code with battery usage in mind
 - Work with, not against the system
 - Pragmatism over principle
 - Efficiency over code elegance and good OO practice
 - Pretend that you're a C programmer

Waking up the phone

- Imagine an email application
 - Checks for new emails every 10 minutes
 - Takes 10 seconds to check for new emails
 - Wakes up CPU using an alarm
 - Wakes up the network device
 - Makes a DNS request
 - Pulls down and parses data
 - 350mA current used
- Cost during a given hour:
 - Sleeping 3600 seconds * 5mA = 5mAh
 - 6 times * 10 seconds * 350mA = 5.8mAh
- Double the battery usage
- Waking up the phone can cause a cascade of updates
 - Some services only run when the phone is awake

Using Alarms

- Multiple scheduled alarms waking the phone
 - Suspending and waking the phone takes power
 - More efficient to schedule multiple alarms at the same time
 - Wake once, do several tasks, sleep once
 - Wake, task, sleep, wake, task, sleep, wake, task, sleep...
- `setInexactRepeating()`
 - An alarm that repeats once an hour, but not necessarily on the hour
 - Time between two alarms may vary
 - Jitter – some may be acceptable
 - Allow the system to schedule multiple similar alarms at the same time

Rules of thumb

- Speed = Efficiency
 - The CPU runs at a certain rate
 - Instructions per second
 - The faster we can perform our work, the more time the CPU can idle
 - Idle at reduced power
 - More efficient use of instructions
 - The faster we can perform our work, the more quickly the CPU can go to sleep
 - Sleeping at reduced power consumption
- Waking up / Running services = Costs power
 - Assume we are not the only application in use
- Byproduct
 - A fast app feels more responsive
 - Users are less likely to use an app that is slow
 - Majority of apps are kept / uninstalled after first run

Memory Use

- Mobiles have limited RAM
 - Typically <512MB
 - iPhone3G used 102MB of 128MB for “system”
 - No virtual memory
 - Be frugal with memory usage
 - Memory allocation / garbage collection takes time, uses power
- Avoid allocation of memory, objects where possible
 - Yes we have a garbage collector
 - Takes time, uses power. Generate less garbage
 - A **String** is an object
 - Return **substring** rather than a new **String** – pointer to the same memory

Throwing away OO

- An array of ints vs an array of Integer objects
 - Which is more efficient in terms of allocation and memory usage?
- Avoid boxing
 - `int x[1024], y[1024];`
 - `class Point { int x; int y; public getX() }`
- More efficient (CPU, memory) to use the former
 - More care required when coding
- Getter / Setter methods
 - Good OO practice to provide these at the class interface
 - Access member variables directly where possible
 - 3x – 7x speed improvement, removing virtual method calls

Throwing away OO

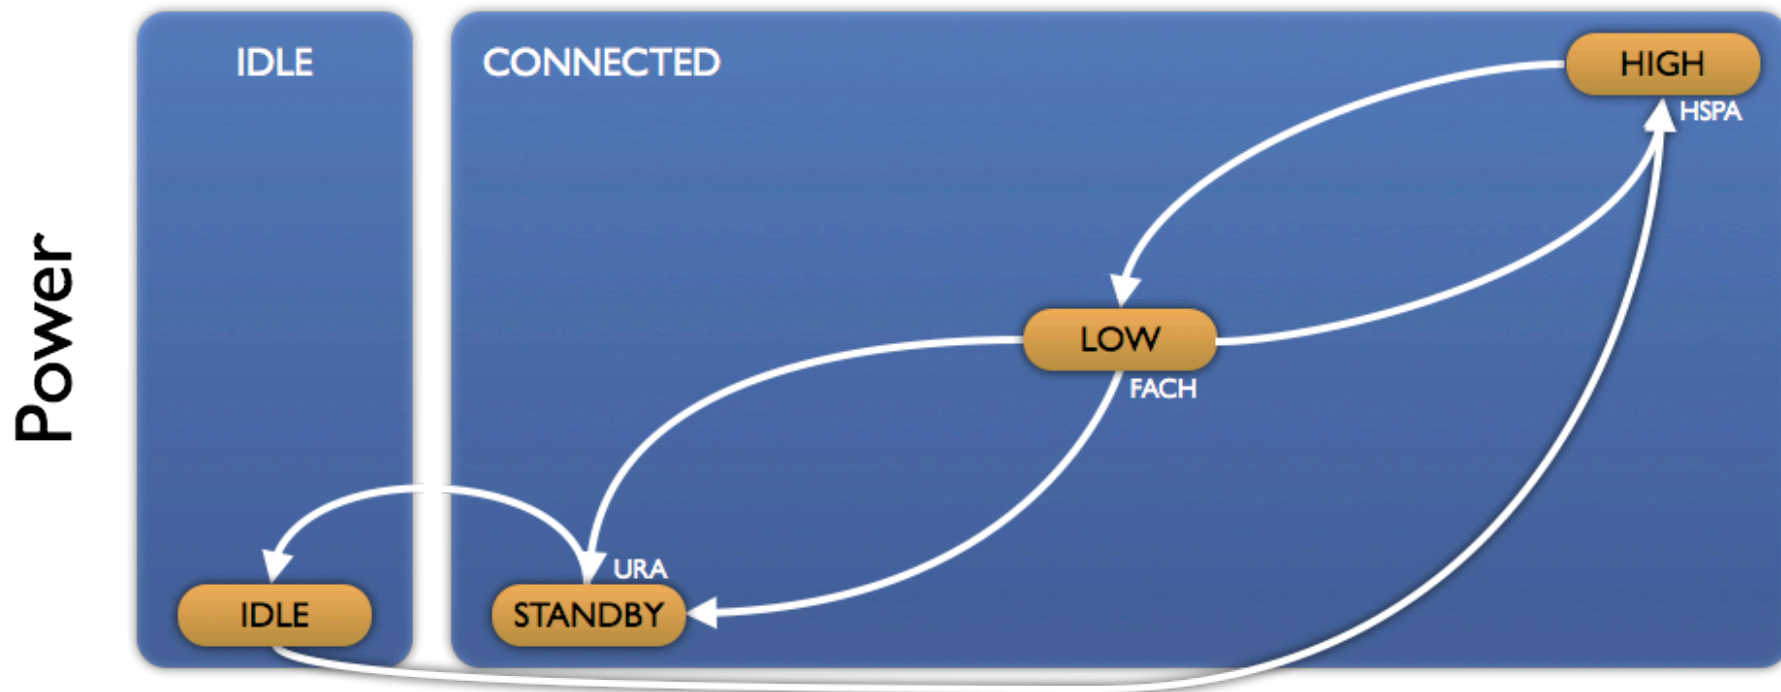
- Static methods belong to a class
 - `public static void foo() { ... }`
- Virtual methods act on an object
 - `public void foo() { ... }`
- When can we use static?
 - When no member variables are accessed
 - Pass variables as parameters instead
 - Makes method invocations ~15% - 20% faster
 - Discuss!
- Constants
 - Should be declared **static final**
 - Just **static** causes class initialiser method to be run
 - Make them **static final** the variable lookup vanishes
 - Replaced by a constant

Other Chips


- CPU ~100mAh
- Other components use power too
 - Accelerometer
 - 10mA normal use – 80mA fastest / finest measurements
 - Choose most appropriate frequency
 - Location
 - Wifi basestations (~100m)
 - Cellphone tower triangulation (~500m – 3km)
 - GPS (~1-5m)
 - Select the most appropriate accuracy
 - GPS is very expensive in terms of battery usage, especially cold start
 - Register for updates appropriately
 - Radios (network connectivity, phone calls)


Radio / Network


- 3G chip has a number of states
 - URA – Connected but not sending data
 - FACH – Half power, small amount of data
 - HSPA – Full power, dedicated channel
- Cost / time to transition upwards
 - Ramp up power, negotiate channel
- In high power radio state
 - Delay to transmit is shorter
 - Device stays in high state for a short period of time following communication
- Regular polling keeps the radio transition between states
 - Pay the battery cost even if we transfer nothing
 - Synchronize polling – inExactAlarms
 - Coalesce data into large chunks
 - Small transfers will only transition up to low / FACH power state (~256 – 512 bytes)
 - Be careful of reusing libraries
 - Were they designed for 3G, or do they assume Ethernet



Data rate / resources / lower latency








Data Transfer

- Battery cost per byte
 - Radio usage, CPU usage
 - Minimise the amount of data transferred
- Reduce signal-to-noise ratio
 - How much of the data describes the structure and not the data?
 - XML is bulkier than JSON
 - JSON is bulkier than binary
- Use Gzip compression where possible
 - Decompressor is native code
 - *Cost to decompress is less than cost to send uncompressed*
- Consider time taken to parse



Source: Timings obtained by downloading and parsing a 6-item RSS feed repeatedly for 60 seconds and averaging results.

References

- <http://www.google.com/events/io/2009/sessions/CodingLifeBatteryLife.html>
- <http://developer.sonymobile.com/2010/08/23/android-tutorial-reducing-power-consumption-of-connected-apps/>
- <http://www.slideshare.net/EricssonLabs/droidcon-understanding-smartphone-traffic>